

- Date.....
- WTWTF/S/SU/ Videos :- DBMS based and Non DBMS Based.
- The Min DBMS
- Big Data Management : The "m" in DBMS.
- DBMS-based and non-DBMS-based Approaches to Big Data.
- Explain the various adv. of using a PBMS over a file system.
- Specify the diff. b/w a parallel and distributed file system.
- MapReduce-style DBMS.
- already in SQL book
- Parallel and distributed DBMS :-
- Parallel database system
    - Improve performance through parallel implementation
    - after allows data replication
    - Data redundancy against table corruption
    - More concurrent queries.
  - Distributed database system:
    - Data is stored across several sites, each site managed by a DBMS capable of running independently!
- \* PBMS and MapReduce-style systems
- Started with a different problem focus
  - DBMSs : efficient storage, transactions and retrieval
    - Partitioned data parallelism
    - Account for computation and communication cost
    - Not node failure
- MapReduce-style systems : complex data processing over a cluster of machines.
- HDFS-based.
- Analytics - data Mining, clustering, me
  - multi-stage, problem - specific algo
  - Operate on wider variety of data including text
    - To climb steep hills requires slow pace at first
- Priority

M/T/W/T/F/S/SU/

Date.....

### Shifting Requirements

- Data loading - a new bottleneck
- Does the app need data sooner than the loading time?
- Too much functionality
- Does the application use only a few data management features?
- Combined Transactional and Analytical capabilities.

### No single solution :-

- Mixed solns
- DBMS on HDFS
  - Hadoop-DBMS interoperation
- Relational operations in MapReduce systems like Spark
- Streaming i/p to DBMS
- New parallel programming models for analytical computation within DBMS.

Priority

- Explain at least 5 desirable characteristics of a Big data management system
- Explain the diff b/w ACID and BASE.
- Describe what the CAP theorem states.
- List examples of BDMS and describe some of their similarities and differences.
- Describe Desired characteristics of BDMS :-
- A flexible, semistructured data model.
  - "schema first" to "schema never"
  - Support for today's common "Big Data data types"
    - Textual, temporal, and spatial data values.
  - A full query language
    - Expectedly at least the power of SQL.
  - An efficient parallel query runtime
  - Wide range of query sizes
  - Continuous data ingestion
    - Stream ingestion
  - Scale gracefully to manage and query large vol<sup>n</sup>s of data
    - Use large clusters.
  - Full data management capability
    - Ease of operational simplicity

ACID and BASE :-

- ACID properties hard to maintain in a BDMS
- BASE relaxes ACID

BA : Basic Availability

S : Soft State

E : Eventual consistency

Priority

M/T/W/T/F/S/SU/

Date.....

CAP Theorem :- A distributed computer system cannot simultaneously achieve

- consistency
- Availability
- Partition Tolerance

Redis :- An Enhanced Key-Value Store

→ In-memory data structure store

- strings, hashes, lists, sets, sorted sets

→ Look-up Problem

- case 1 : (key : string, value : string)

→ keys may have internal structure and expiry

- case 2 : (key : string, value : list)

→ Ziplists compress lists

→ Twitter innovation : list of ziplists

- case 3 : (key : string, value : attribute-value pairs)

→ REDIS Hashes

# Redis and scalability

- Partitioning and Replication

① Range Partitioning :-

Example :- User record number 1-10000 get to machine 1

— 11 —

10001 - 20000 get to machine 2

② Hash partitioning

→ pick a key of a record eg. "abcde"

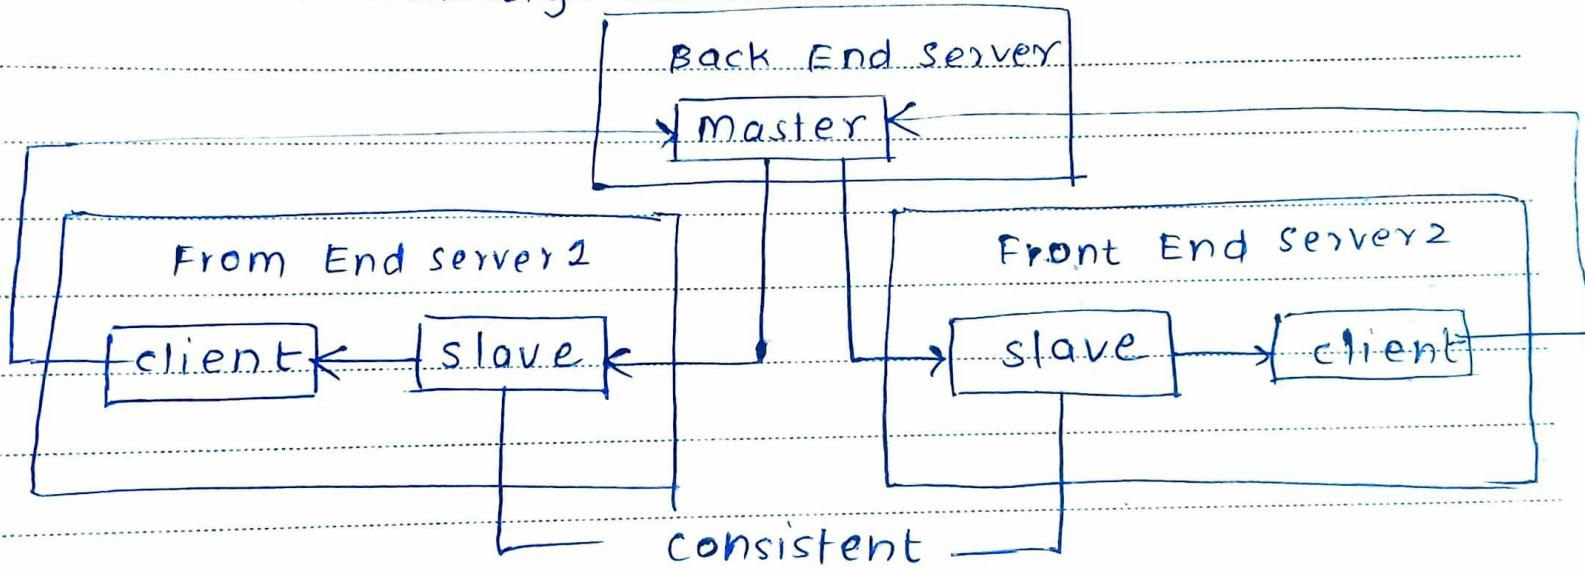
→ Using a hash function, turn it into a no.

→  $152 \bmod 10$  is 2, so the record goes to machine 2

Priority

- To avoid criticism say nothing, do nothing, be nothing

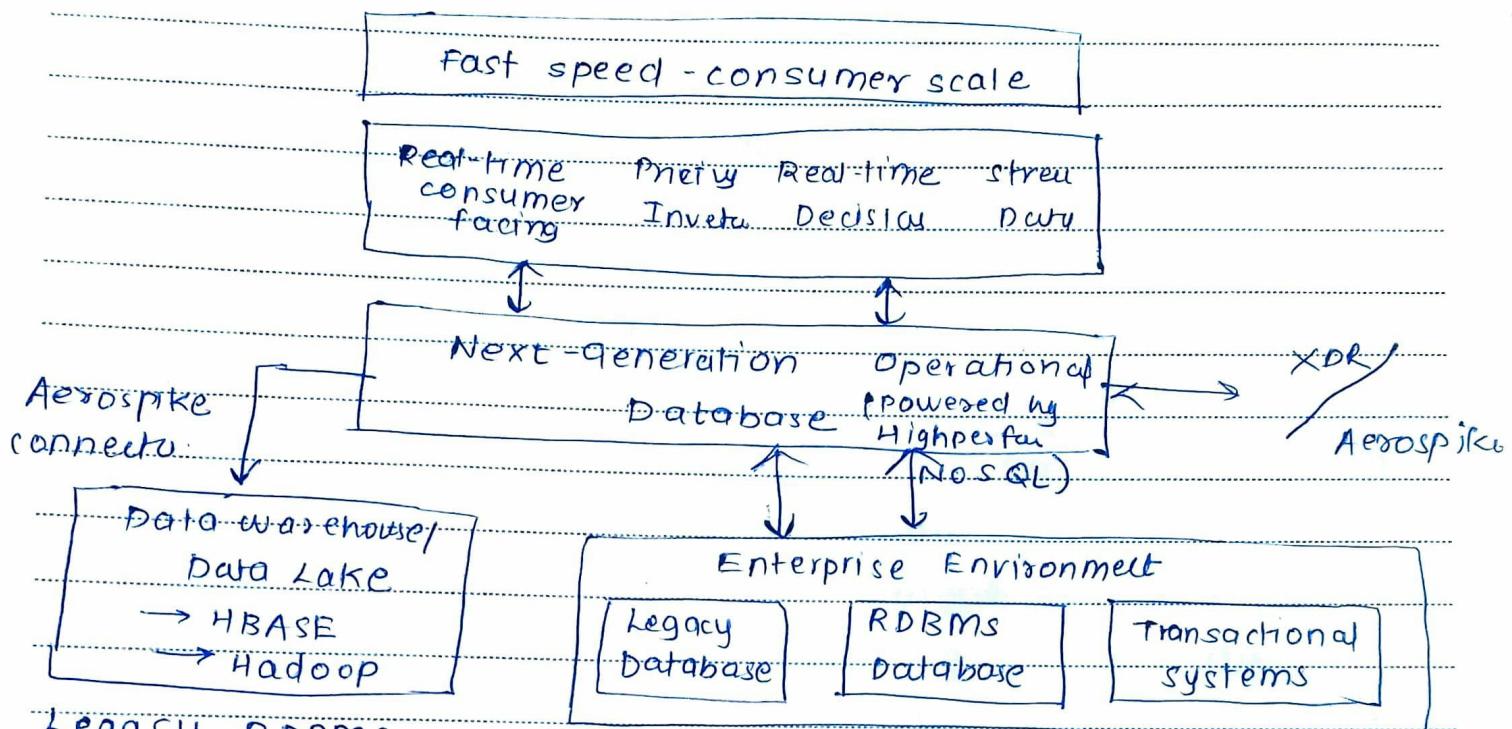
- Partitioning and Replication:
- Master-slave mode replication
- client write to master, master replicates to slaves.
  - clients read from slaves to scale up read performance.
  - slaves are mostly consistent.



M/T/W/T/F/S/SU/

Date.....

Aerospike :- a New Generation KV Store

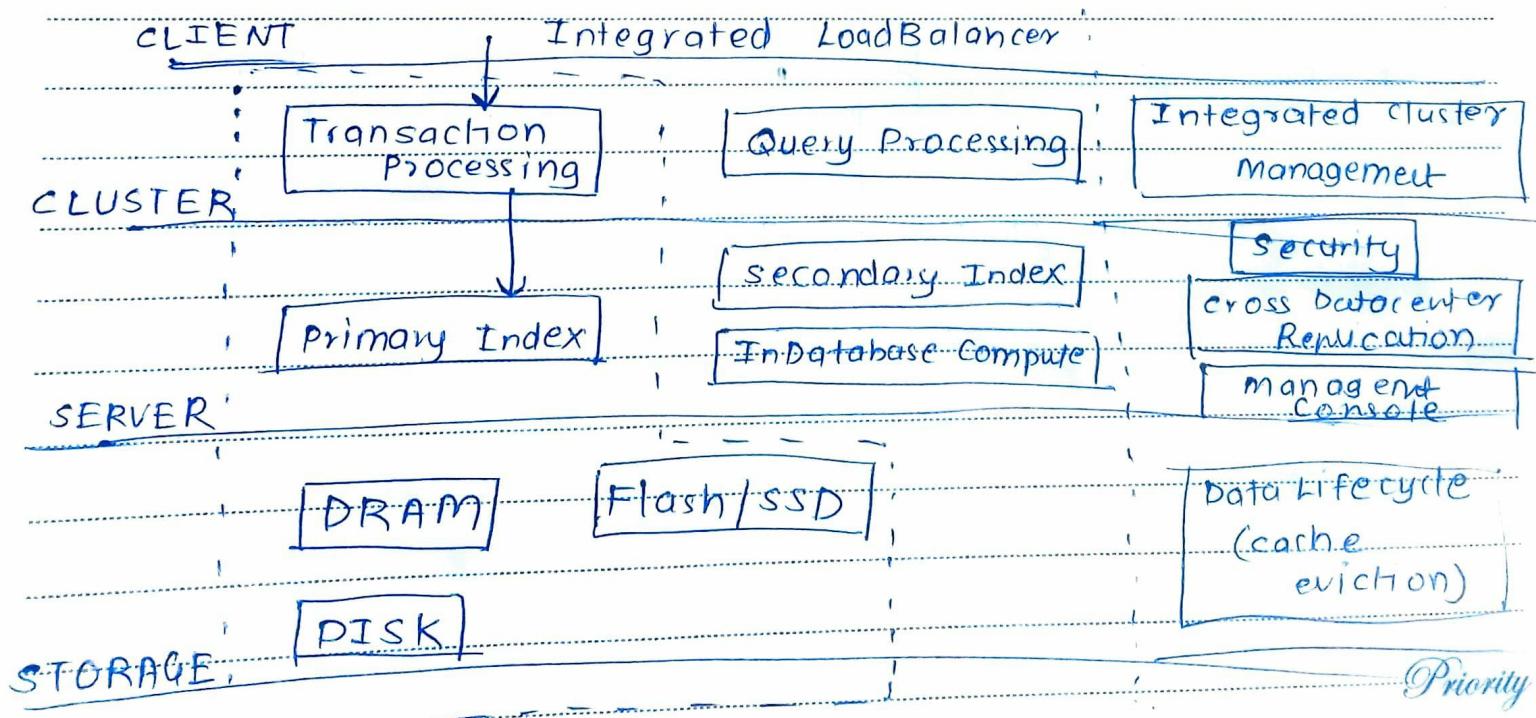


Legacy RDBMS

HDFS Based

\* Aerospike Architecture :-

Key Value store (Fast Path)      Operational Queries      Management



- To enjoy life one should give up the lure of life

August 2016						
Wk	M	T	W	T	F	S
32	1	2	3	4	5	6
33	8	9	10	11	12	13
34	15	16	17	18	19	20
35	22	23	24	25	26	27
36	29	30	31			

September 2016						
Wk	M	T	W	T	F	S
36				1	2	3
37		5	6	7	8	9
38		12	13	14	15	16
39		19	20	21	22	23
40		26	27	28	29	30

2016  
Thursday  
July

07

189-177 • WK 28

## Querying Aerospike :-

- Data types :- standard scalar, lists, maps, geospatial, large objects.
- KV store operations :- geospatial queries like point-in-polygon.
- AQL : an SQL-like language.

SELECT name, age

FROM users.profiles

## Transactions in Aerospike :-

Aerospike ensures ACID

- consistency :- all copies of a data item are in sync
- uses synchronous write to replicas
- Mechanisms to relax immediate consistency

Durability :- → Flash storage

→ Replication management

Network partitioning reduced

→ Tighter cluster control

5.00

AsterixDB :- a DBMS for semistructured data

## # Options for Querying in AsterixDB

- AQL is a natively-supported query language
- for \$user in dataset TwitterUsers
- order by \$user.followers\_count desc,
- \$user.lang asc return \$user.

## NOTES Hive queries

SELECT a-val, b-val

FROM a LEFT OUTER JOIN b ON (a.key=b.key)

→ Xquery

→ Hadoop MR jobs

→ SQL ft (coming up )

June 2016							July 2016										
06		M	T	W	T	F	S	S	07	wk	M	T	W	T	F	S	S
wk									wk								
23					1	2	3	4	27					1	2	3	
24	6	7	8	9	10	11	12		28	4	5	6	7	8	9	10	
25	13	14	15	16	17	18	19		29	11	12	13	14	15	16	17	
26	20	21	22	23	24	25	26		30	18	19	20	21	22	23	24	
27	27	28	29	30					31	25	26	27	28	29	30	31	

\* Operating over a cluster.

- 9.00 • Hyracks :- Query execution engine for partitioned, parallel execution of queries.

Ex :- CUSTOMER ( C-CUSTKEY, C-MKTSEGMENT, ... )  
ORDERS ( O-ORDERKEY, O-CUSTKEY, ... )

11.00

## Accessing External Data :-

- Real-time data from files in a directory path.
  - Real-time data from an external API.

100

## Solr :- Managing Text

- Basic challenges with text
    - Defining a match

Analyze = Analyse

Dr = Doctor

## 4.00 Inverted Index

- Vocabulary :- All terms in a collection of documents
    - multi-word terms, synonym sets.
  - Occurrence :- For each term in the collection
    - List of docID
    - List of docID, [position of occurrence]

## ~~Solar Functionality :-~~

- Enterprise search Platform
  - Inverted index
    - For every field in a structured text document

## NOTES

→ Indexes text, numbers

## → Faceted Search

## → term highlighting

→ Solr Functionality :- • Index-time Analyzers - Tokenizers  
- Filter

Wk	M	T	W	T	F	S	S
36			1	2	3	4	
37	5	6	7	8	9	10	11
38	12	13	14	15	16	17	18
39	19	20	21	22	23	24	25
40	26	27	28	29	30		

Wk	M	T	W	T	F	S	S
40/45			31				
41	3	4	5	6	7	8	9
42	10	11	12	13	14	15	16
43	17	18	19	20	21	22	23
44	24	25	26	27	28	29	30

2016  
Wednesday  
August

10

223-143 • WK 33

Vertical :- a columnar DBMS

Row-oriented database

column-oriented database

- column store
- store data column-wise
- A query only uses the columns needed.
- Usually much faster for queries even for large data.

### \* Space Efficiency

→ Column stores keep column in sorted order.

→ Values in columns can be compressed.

- Run-length encoding
- Frame-of-references encoding

→ Compression saves storage space.

### Working with Vertical :-

• Column-groups

→ Frequently co-accessed columns behave as mini-row-stores within the column store.

• Update performance slower.

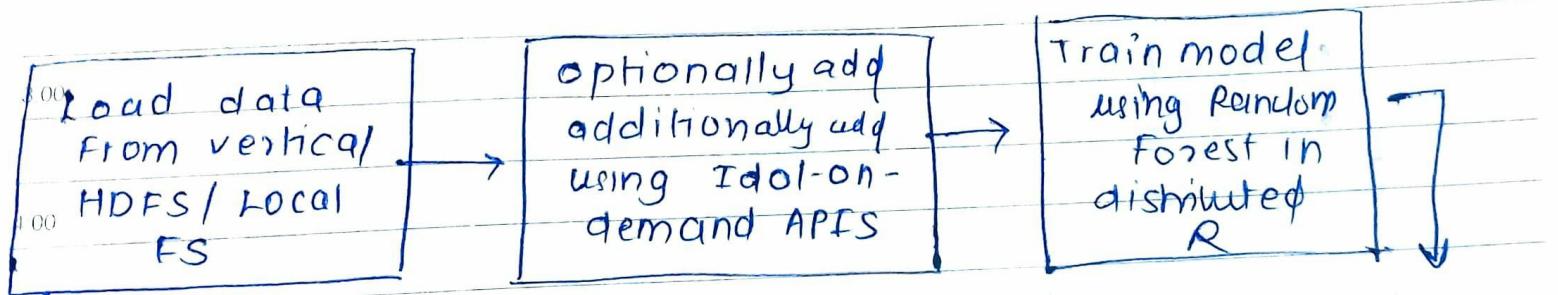
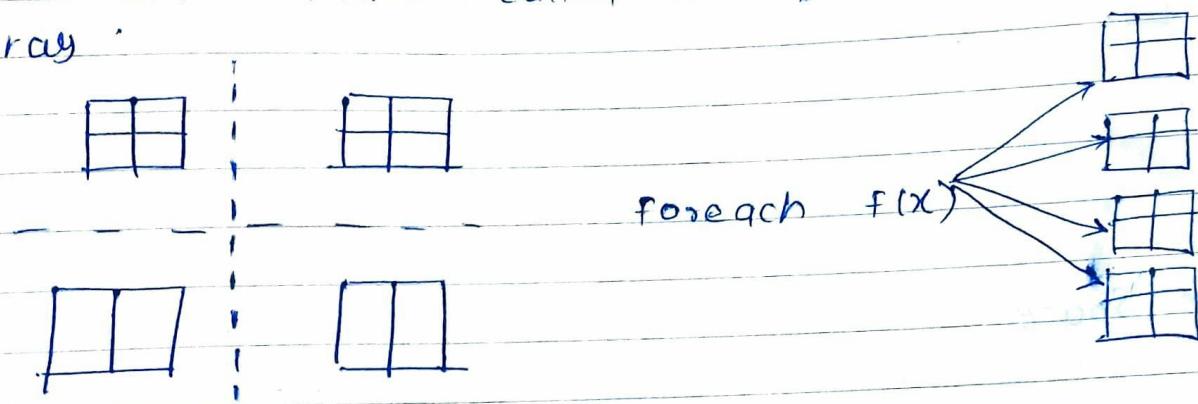
→ Internal conversion from row-representation to column-representation.

• Enhanced suite of analytical operations

DIES → Windows statistics.

## Vertical and Distributed R :-

- High-performance statistical analysis.
- Master node : schedules tasks and sends code.
- Worker node : maintain data partitions and compute.
- Uses a data structure called darray or distributed array.



NOTES